

INDEX

SR.NO	PRACTICAL NAME	PRACTICAL DATE	TEACHER'S SIGNATURE
1.	Document Indexing and Retrieval		
2.	Retrieval Models		
3.	Spelling Correction in IR Systems		
4.	Evaluation Metrics for IR Systems		
5.	Text Categorization		
6.	Clustering for Information Retrieval		
7.	Web Crawling and Indexing		
8.	Link Analysis and PageRank		
9.	Learning to Rank		
10.	Advanced Topics in Information Retrieval		

Practical No.1

Aim:- Document Indexing and Retrieval

- **Implement an inverted index construction algorithm.**

Code:-

```
document1 = "The quick brown fox jumped over the lazy  
dog."
```

```
document2 = "The lazy dog slept in the sun."
```

```
tokens1 = document1.lower().split()
```

```
tokens2 = document2.lower().split()
```

```
terms = list(set(tokens1 + tokens2))
```

```
inverted_index = {}
```

```
for term in terms:
```

```
    documents = []
```

```
    if term in tokens1:
```

```
        documents.append("Documents 1")
```

```
    if term in tokens2:
```

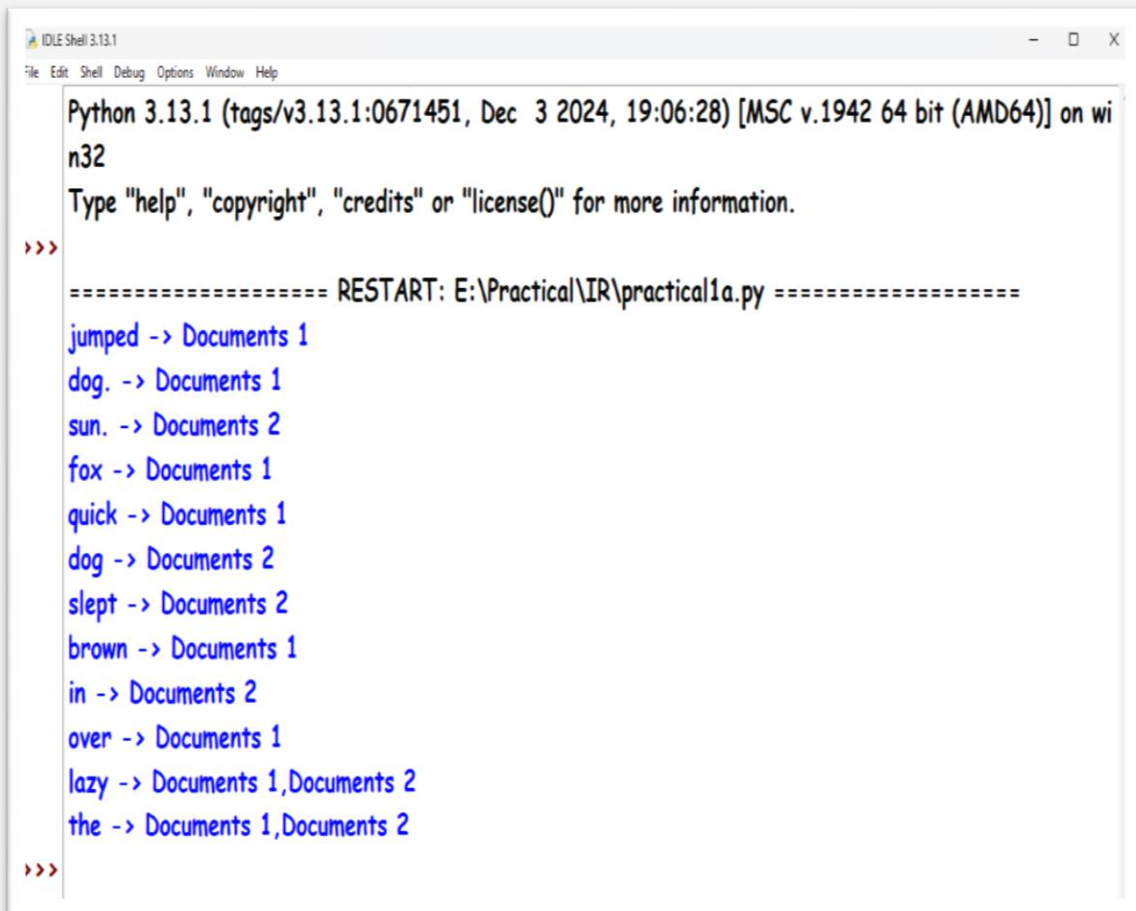
```
documents.append("Documents 2")
```

```
inverted_index[term] = documents
```

```
for term, documents in inverted_index.items():
```

```
    print(term, "->", ",".join(documents))
```

Output:-

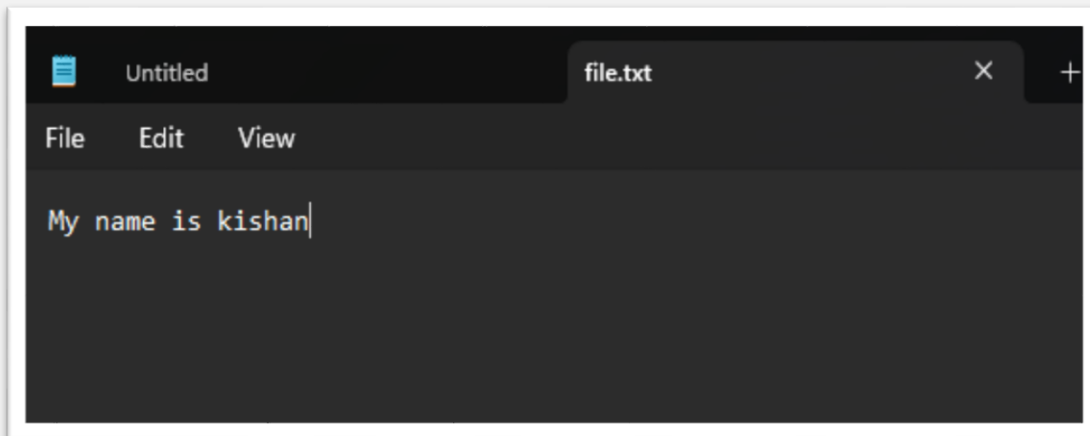


```
IDLE Shell 3.13.1
File Edit Shell Debug Options Window Help
Python 3.13.1 (tags/v3.13.1:0671451, Dec 3 2024, 19:06:28) [MSC v.1942 64 bit (AMD64)] on win32
Type "help", "copyright", "credits" or "license()" for more information.
>>>
===== RESTART: E:\Practical\IR\practical1a.py =====
jumped -> Documents 1
dog. -> Documents 1
sun. -> Documents 2
fox -> Documents 1
quick -> Documents 1
dog -> Documents 2
slept -> Documents 2
brown -> Documents 1
in -> Documents 2
over -> Documents 1
lazy -> Documents 1, Documents 2
the -> Documents 1, Documents 2
>>>
```

- **Build a simple document retrieval system using the constructed index.**

1 B)

Step 1:- Create the text file and file.txt



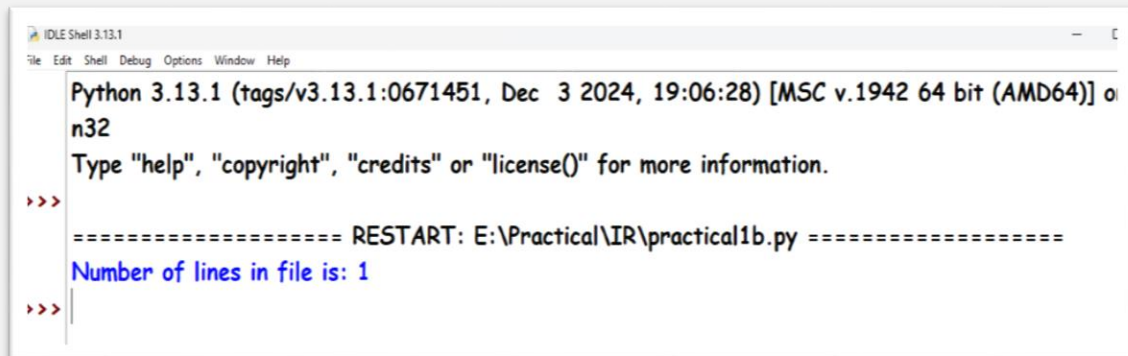
Step 2:- Write python code

Code:-

```
file=open('file.txt',encoding='utf8')
read = file.read()
file.seek(0)
read
line = 1
for word in read:
    if word == '\n':
        line += 1
print("Number of lines in file is:",line)
```

```
array=[]  
for i in range(line):  
    array.append(file.readline())
```

Output:-

A screenshot of an IDLE Shell window for Python 3.13.1. The window title is 'IDLE Shell 3.13.1'. The menu bar includes 'File', 'Edit', 'Shell', 'Debug', 'Options', 'Window', and 'Help'. The shell output shows the Python version and build information: 'Python 3.13.1 (tags/v3.13.1:0671451, Dec 3 2024, 19:06:28) [MSC v.1942 64 bit (AMD64)] on n32'. It then prompts the user to type 'help', 'copyright', 'credits' or 'license()' for more information. Below this, there is a restart message: '==== RESTART: E:\Practical\IR\practical1b.py ====='. The output of the program is displayed in blue text: 'Number of lines in file is: 1'. The prompt '»»»' is visible at the end of the line.

1C)

Code:-

```
from nltk.tokenize import word_tokenize
import nltk
from nltk.corpus import stopwords
nltk.download('stopwords')
nltk.download('punkt')
nltk.download('punkt_tab')
```

```
read = "This is a simple sentence"
```

```
text_tokens = word_tokenize(read)
```

```
tokens_without_sw = [word for word in text_tokens if
word.lower() not in stopwords.words('english')]
print(tokens_without_sw)
```

Output:-

```
IDLE Shell 3.13.1
File Edit Shell Debug Options Window Help
Python 3.13.1 (tags/v3.13.1:0671451, Dec 3 2024, 19:06:28) [MSC v.1942 64 bit (AMD64)] on n32
Type "help", "copyright", "credits" or "license()" for more information.
>>>
===== RESTART: E:\Practical\IR\practical1c.py =====
[nltk_data] Downloading package stopwords to
[nltk_data]   C:\Users\kc140\AppData\Roaming\nltk_data...
[nltk_data] Package stopwords is already up-to-date!
[nltk_data] Downloading package punkt to
[nltk_data]   C:\Users\kc140\AppData\Roaming\nltk_data...
[nltk_data] Package punkt is already up-to-date!
[nltk_data] Downloading package punkt_tab to
[nltk_data]   C:\Users\kc140\AppData\Roaming\nltk_data...
[nltk_data] Unzipping tokenizers\punkt_tab.zip.
['simple', 'sentence']
```

Practical No.2

Aim:- Retrieval Models

- Implement the Boolean retrieval model and process queries.

Code:-

```
from collections import defaultdict
import re

documents = {
    1: "The quick brown fox jumps over the lazy dog.",
    2: "Never jump over the lazy dog quickly.",
    3: "Foxes are quick and smart animals."
}

def preprocess(text):
    text = text.lower()
    text = re.sub(r'\W+', ' ', text)
    return text.split()

def build_inverted_index(docs):
    inverted_index = defaultdict(list)
    for doc_id, text in docs.items():
        words = preprocess(text)
        for word in set(words):
```



```
        inverted_index[word].append(doc_id)
    return inverted_index

inverted_index = build_inverted_index(documents)

def search(query, inverted_index):
    query_words = preprocess(query)
    result_sets = [ ]

    for word in query_words:
        result_sets.append(set(inverted_index.get(word, [ ])))

    if result_sets:
        return set.intersection(*result_sets)
    else:
        return set( )

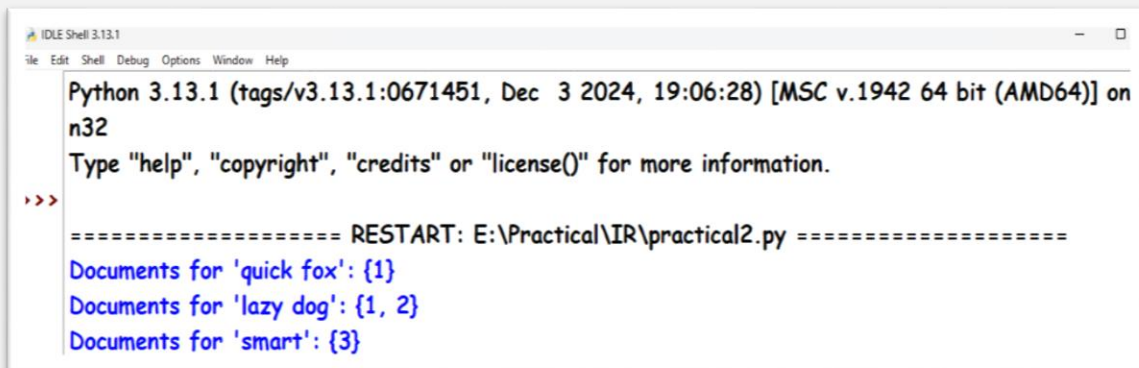
query1 = "quick fox"
query2 = "lazy dog"
query3 = "smart"
```

```
print(f'Documents for '{query1}': {search(query1,  
inverted_index)}")
```

```
print(f'Documents for '{query2}': {search(query2,  
inverted_index)}")
```

```
print(f'Documents for '{query3}': {search(query3,  
inverted_index)}")
```

Output:-



```
IDLE Shell 3.13.1
File Edit Shell Debug Options Window Help
Python 3.13.1 (tags/v3.13.1:0671451, Dec 3 2024, 19:06:28) [MSC v.1942 64 bit (AMD64)] on
n32
Type "help", "copyright", "credits" or "license()" for more information.
»»»
===== RESTART: E:\Practical\IR\practical2.py =====
Documents for 'quick fox': {1}
Documents for 'lazy dog': {1, 2}
Documents for 'smart': {3}
```

- **Implement the vector space model with TF-IDF weighting and cosine similarity.**

Code:-

```
from sklearn.feature_extraction.text import CountVectorizer,
TfidfTransformer

import nltk

from nltk.corpus import stopwords

import numpy as np

from numpy.linalg import norm

train_set = ["The sky is blue.", "The sun is bright."] #
Documents

test_set = ["The sun in the sky is bright."] # Query

nltk.download('stopwords')

stopWords = stopwords.words('english')

vectorizer = CountVectorizer(stop_words=stopWords)

transformer = TfidfTransformer()

trainVectorizerArray =
vectorizer.fit_transform(train_set).toarray()

testVectorizerArray = vectorizer.transform(test_set).toarray()

print('Fit Vectorizer to train set:', trainVectorizerArray)

print('Transform Vectorizer to test set:', testVectorizerArray)
```

```
cx = lambda a, b: round(np.inner(a, b) / (norm(a) * norm(b)),  
3)
```

```
for vector in trainVectorizerArray:
```

```
    print(vector)
```

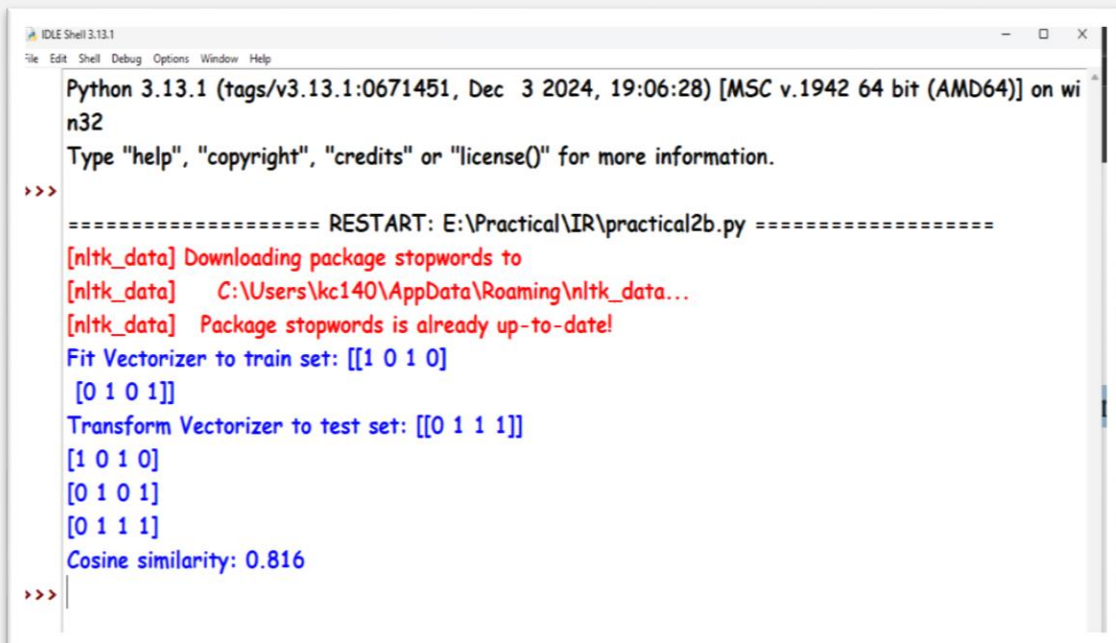
```
for testV in testVectorizerArray:
```

```
    print(testV)
```

```
    cosine = cx(vector, testV)
```

```
print(f"Cosine similarity: {cosine}")
```

Output:-



```
IDLE Shell 3.13.1
Python 3.13.1 (tags/v3.13.1:0671451, Dec 3 2024, 19:06:28) [MSC v.1942 64 bit (AMD64)] on win32
Type "help", "copyright", "credits" or "license()" for more information.
>>>
===== RESTART: E:\Practical\IR\practical2b.py =====
[nltk_data] Downloading package stopwords to
[nltk_data] C:\Users\kc140\AppData\Roaming\nltk_data...
[nltk_data] Package stopwords is already up-to-date!
Fit Vectorizer to train set: [[1 0 1 0]
[0 1 0 1]]
Transform Vectorizer to test set: [[0 1 1 1]]
[1 0 1 0]
[0 1 0 1]
[0 1 1 1]
Cosine similarity: 0.816
>>>
```

Practical No.3

Aim:- Spelling Correction in IR Systems

- **Develop a spelling correction module using edit distance algorithms.**
- **Integrate the spelling correction module into an information retrieval system**

Code:-

```
def editDistance(str1,str2, m, n):  
    if m == 0:  
        return n  
    if n == 0:  
        return m  
    if str1[m-1] == str2[n-1]:  
        return editDistance(str1, str2, m-1, n-1)  
    return 1 + min(  
        editDistance(str1, str2, m, n-1),  
        editDistance(str1, str2, m-1, n),  
        editDistance(str1, str2, m-1, n-1)  
    )  
str1 = "sunday"  
str2 = "saturday"
```

```
print('Edit Distance is:', editDistance(str1, str2, len(str1),  
len(str2)))
```

Output:-

```
===== RESTART: E:\Practical\LR\Practical 3.py =====  
Edit Distance is: 3  
>>>
```

Practical No.4

Aim:- Evaluation Metrics for IR Systems.

- **Calculate precision, recall, and F-measure for a given set of retrieval results.**

Code:-

```
def calculate_precision_recall_fmeasure(relevant_docs,
retrieved_docs):
```

```
    relevant_set = set(relevant_docs)
```

```
    retrieved_set = set(retrieved_docs)
```

```
    relevant_retrieved = len(relevant_set & retrieved_set)
```

```
    precision = relevant_retrieved / len(retrieved_set) if
len(retrieved_set) > 0 else 0
```

```
    recall = relevant_retrieved / len(relevant_set) if
len(relevant_set) > 0 else 0
```

```
    if precision + recall > 0:
```

```
        f_measure = 2 * (precision * recall) / (precision +
recall)
```

```
    else:
```

```
        f_measure = 0
```

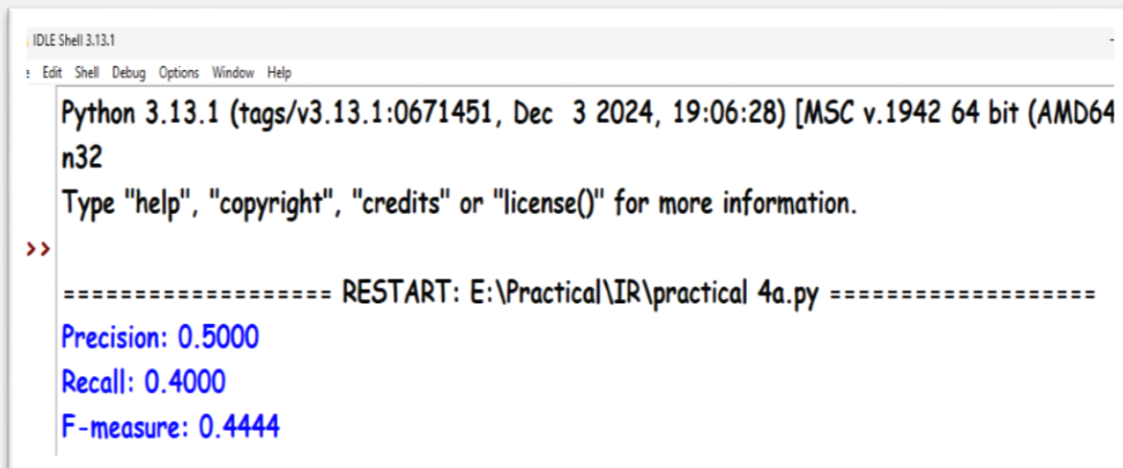
```
    return precision, recall, f_measure
```

```
relevant_docs = [1, 2, 3, 4, 5]
```

```
retrieved_docs = [2, 4, 6, 7]
```

```
precision, recall, f_measure =  
calculate_precision_recall_fmeasure(relevant_docs,  
retrieved_docs)  
  
print(f"Precision: {precision:.4f}")  
print(f"Recall: {recall:.4f}")  
print(f"F-measure: {f_measure:.4f}")
```

Output:-



```
IDLE Shell 3.13.1  
File Edit Shell Debug Options Window Help  
Python 3.13.1 (tags/v3.13.1:0671451, Dec 3 2024, 19:06:28) [MSC v.1942 64 bit (AMD64)]  
n32  
Type "help", "copyright", "credits" or "license()" for more information.  
>>  
===== RESTART: E:\Practical\IR\practical 4a.py =====  
Precision: 0.5000  
Recall: 0.4000  
F-measure: 0.4444
```


- **Use an evaluation toolkit to measure average precision and other evaluation metrics.**

Code:-

```
from sklearn.metrics import average_precision_score
y_true = [0, 1, 1, 0, 1, 1]
y_scores = [0.1, 0.4, 0.35, 0.8, 0.65, 0.9]
average_precision = average_precision_score(y_true,
y_scores)
print(f'Average precision-recall score: {average_precision}')
```

Output:-

```
===== RESTART: E:\Practical\IR\practical 4.py =====
Average precision-recall score: 0.8041666666666667
>>|
```

Practical No.5

Aim :- Text Categorization

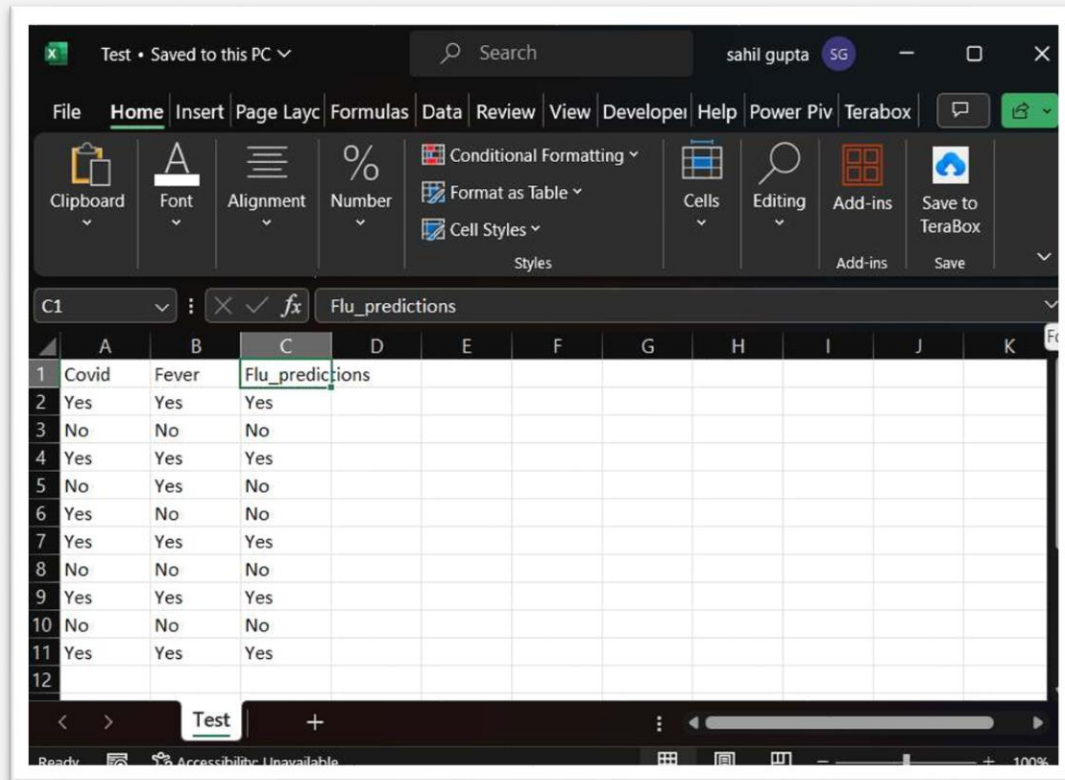
- Implement a text classification algorithm (e.g., Naive Bayes or Support Vector Machines).
- Train the classifier on a labelled dataset and evaluate its performance.

Step 1:- Create a two csv file first is Dataset.csv and second is Test.csv

Dataset.csv

	A	B	C	D	E	F	G	H	I	J
1	Covid	Fever	Flu							
2	Yes	Yes	Yes							
3	No	No	No							
4	No	Yes	No							
5	No	Yes	No							
6	Yes	No	No							
7	Yes	Yes	Yes							
8	No	No	No							
9	Yes	Yes	Yes							
10	No	No	No							
11	No	Yes	No							
12										

Test.csv



Code:-

```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.feature_extraction.text import
CountVectorizer
from sklearn.naive_bayes import MultinomialNB
from sklearn.metrics import accuracy_score,
classification_report
df =
pd.read_csv(r"C:\Users\kc140\OneDrive\Documents\Datas
et.csv")
data = df["Covid"] + " " + df["Fever"]
```

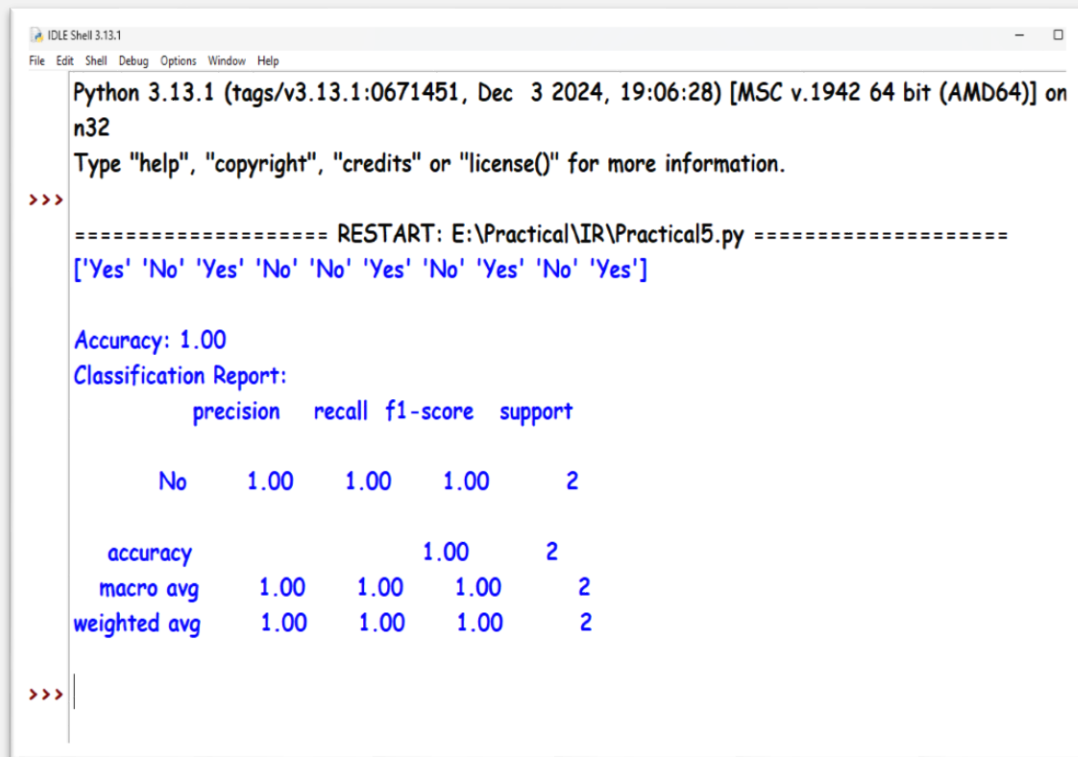
```

X = data.astype(str)
y = df['Flu']
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)
vectorizer = CountVectorizer()
X_train_counts = vectorizer.fit_transform(X_train)
X_test_counts = vectorizer.transform(X_test)
classifier = MultinomialNB()
classifier.fit(X_train_counts, y_train)
data1 =
pd.read_csv(r"C:\Users\kc140\OneDrive\Documents\Test.c
sv")
new_data = data1["Covid"] + " " + data1["Fever"]
new_data_counts =
vectorizer.transform(new_data.astype(str))
predictions = classifier.predict(new_data_counts)
new_data = predictions
print(new_data)
accuracy = accuracy_score(y_test,
classifier.predict(X_test_counts))
print(f"\nAccuracy: {accuracy:.2f}")
print("Classification Report:")

```

```
print(classification_report(y_test,  
classifier.predict(X_test_counts)))
```

Output:-



```
IDLE Shell 3.13.1
Python 3.13.1 (tags/v3.13.1:0671451, Dec 3 2024, 19:06:28) [MSC v.1942 64 bit (AMD64)] on
n32
Type "help", "copyright", "credits" or "license()" for more information.
>>>
===== RESTART: E:\Practical\IR\Practical5.py =====
['Yes' 'No' 'Yes' 'No' 'No' 'Yes' 'No' 'Yes' 'No' 'Yes']

Accuracy: 1.00
Classification Report:
      precision    recall  f1-score   support

      No         1.00      1.00      1.00         2

 accuracy                   1.00         2
 macro avg           1.00      1.00      1.00         2
 weighted avg        1.00      1.00      1.00         2

>>>
```

Practical No.6

Aim:- Clustering for Information Retrieval

- **Implement a clustering algorithm (e.g., K-means or hierarchical clustering).**
- **Apply the clustering algorithm to a set of documents and evaluate the clustering results.**

Code:-

```
from sklearn.feature_extraction.text import  
TfidfVectorizer  
  
from sklearn.cluster import KMeans  
  
documents = ["Cats are known for their agility and  
grace",  
"Dogs are often called ‘man’s best friend’.",  
"Some dogs are trained to assist people with  
disabilities.",  
"The sun rises in the east and sets in the west.",  
"Many cats enjoy climbing trees and chasing toys.",  
]  
  
vectorizer = TfidfVectorizer(stop_words='english')  
  
X = vectorizer.fit_transform(documents)  
  
kmeans = KMeans(n_clusters=3, random_state=0).fit(X)  
  
print(kmeans.labels_)
```

Output:-

```
===== RESTART: E:\Practical\IR\Practical6.py =====  
[0 2 0 1 0]  
>> |
```

Practical No.7

Aim:- Web Crawling and Indexing.

- **Develop a web crawler to fetch and index web pages.**
- **Handle challenges such as robots.txt, dynamic content, and crawling delays.**

Code:-

```
import requests
from bs4 import BeautifulSoup
import time
from urllib.parse import urljoin, urlparse
from urllib.robotparser import RobotFileParser

def get_html(url):
    headers = {'User-Agent': 'Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/58.0.3029.110 Safari/537.3'}
    try:
        response = requests.get(url, headers=headers)
        response.raise_for_status()
        return response.text
    except requests.exceptions.HTTPError as errh:
        print(f"HTTP Error: {errh}")
```



```
except requests.exceptions.RequestException as err:
    print(f'Request Error: {err}')
return None
```

```
def save_robots_txt(url):
    try:
        robots_url = urljoin(url, '/robots.txt')
        robots_content = get_html(robots_url)
        if robots_content:
            with open('robots.txt', 'wb') as file:
                file.write(robots_content.encode('utf-8-sig'))
    except Exception as e:
        print(f'Error saving robots.txt: {e}')
```

```
def load_robots_txt():
    try:
        with open('robots.txt', 'rb') as file:
            return file.read().decode('utf-8-sig')
    except FileNotFoundError:
        return None
```

```

def extract_links(html, base_url):
    soup = BeautifulSoup(html, 'html.parser')
    links = []
    for link in soup.find_all('a', href=True):
        absolute_url = urljoin(base_url, link.get('href'))
        links.append(absolute_url)
    return links

def is_allowed_by_robots(url, robots_content):
    parser = RobotFileParser()
    parser.parse(robots_content.split('\n'))
    return parser.can_fetch('*', url)

def crawl(start_url, max_depth=3, delay=1):
    visited_urls = set()

    def recursive_crawl(url, depth, robots_content):
        if depth > max_depth or url in visited_urls or not
is_allowed_by_robots(url, robots_content):
            return
        visited_urls.add(url)
        time.sleep(delay)
        html = get_html(url)
        if html:
            print(f'Crawling {url}')
            links = extract_links(html, url)
            for link in links:
                recursive_crawl(link, depth + 1, robots_content)

    save_robots_txt(start_url)
    robots_content = load_robots_txt()

```

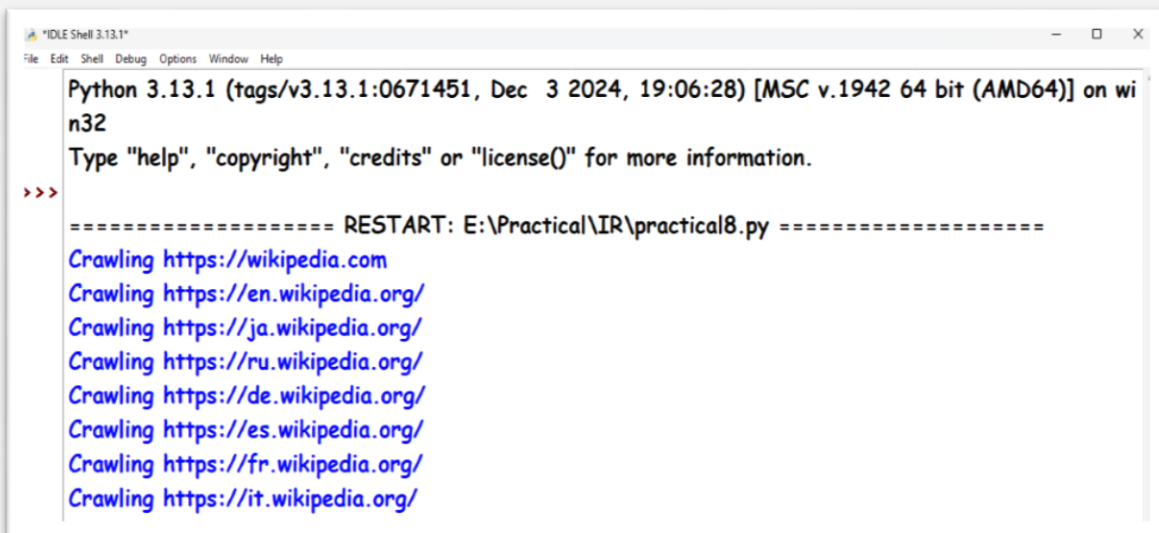
```
if not robots_content:  
    print("Unable to retrieve robots.txt. Crawling without  
restrictions.")
```

```
recursive_crawl(start_url, 1, robots_content)  
print("Performed by 740_Pallavi & 743_Deepak")
```

Start crawling from Wikipedia with max depth 2 and delay of 2 seconds

```
crawl('https://wikipedia.com', max_depth=2, delay=2)
```

Output :-



```
"IDLE Shell 3.13.1"  
File Edit Shell Debug Options Window Help  
Python 3.13.1 (tags/v3.13.1:0671451, Dec 3 2024, 19:06:28) [MSC v.1942 64 bit (AMD64)] on wi  
n32  
Type "help", "copyright", "credits" or "license()" for more information.  
>>>  
===== RESTART: E:\Practical\IR\practical8.py =====  
Crawling https://wikipedia.com  
Crawling https://en.wikipedia.org/  
Crawling https://ja.wikipedia.org/  
Crawling https://ru.wikipedia.org/  
Crawling https://de.wikipedia.org/  
Crawling https://es.wikipedia.org/  
Crawling https://fr.wikipedia.org/  
Crawling https://it.wikipedia.org/
```

Practical No.8

Aim :- Link Analysis and PageRank.

- **Implement the PageRank algorithm to rank web pages based on link analysis.**

Code:-

```
import numpy as np

def pagerank(graph, d=0.85, max_iter=100, tol=1e-6):
    n = len(graph)
    ranks = np.ones(n) / n
    M = np.array(graph, dtype=float)
    outdegree = M.sum(axis=0)
    M = np.divide(M, outdegree, where=outdegree != 0)
    for _ in range(max_iter):
        new_ranks = (1 - d) / n + d * np.dot(M, ranks)
        if np.linalg.norm(new_ranks - ranks, 1) < tol:
            break
        ranks = new_ranks
    return ranks

graph = [
    [0, 1, 1, 0],
    [0, 0, 1, 1],
```

```
[1, 0, 0, 1],  
[1, 0, 0, 0],  
]  
ranks = pagerank(graph)  
print("PageRank of each page:", ranks)
```

Output:-

```
type "help", "copyright", "credits" or "license()" for more information.  
>>>  
===== RESTART: E:\Practical\IR\pagerank_algo.py =====  
PageRank of each page: [0.33739791 0.22393378 0.25777409 0.18089421]  
>>>
```

- **Apply the PageRank algorithm to a small web graph and analyze the results.**

Code:-

```
import networkx as nx

G = nx.DiGraph()

G.add_edges_from([('A', 'B'), ('A', 'C'), ('B', 'C'), ('C', 'A'), ('C', 'B')])

pagerank_scores = nx.pagerank(G, alpha=0.85)

for node, score in pagerank_scores.items():
    print(f'Node {node} has a PageRank score of {score:.4f}')
```

Output :-

```
>>> ..... RESTART: E:/Practical/IR/practical8b.py .....
Node A has a PageRank score of 0.2339
Node B has a PageRank score of 0.3333
Node C has a PageRank score of 0.4327
```

Practical No.9

Aim:- Learning to Rank

- **Implement a learning to rank algorithm (e.g. RankSVM or RankBoost).**

Code:-

```
import numpy as np
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error

queries = [
    "best programming languages",
    "machine learning tutorials",
    "artificial intelligence books"
]

documents = [
    "Python is a versatile language used for machine learning and web development.",
    "This guide teaches machine learning using scikit-learn and TensorFlow.",
    "AI books are essential for learning deep learning techniques and neural networks.",
```

"Java is widely used for enterprise applications and Android development.",

"Books on AI cover topics from basic to advanced deep learning algorithms."

]

relevance_scores = [

[3, 5, 2, 1, 4],

[4, 5, 3, 2, 1],

[5, 4, 3, 1, 2]

]

vectorizer = TfidfVectorizer(stop_words='english')

X_documents = vectorizer.fit_transform(documents)

y = np.array(relevance_scores).flatten()

X_query_doc = []

for query in queries:

 query_vector = vectorizer.transform([query]) # Vectorize
the query

 query_doc_pairs =
[np.concatenate([query_vector.toarray()[0], doc.toarray()[0]])
for doc in X_documents]

 X_query_doc.extend(query_doc_pairs)


```
X_query_doc = np.array(X_query_doc)
```

```
X_train, X_test, y_train, y_test =  
train_test_split(X_query_doc, y, test_size=0.2,  
random_state=42)
```

```
model = LinearRegression()
```

```
model.fit(X_train, y_train)
```

```
y_pred = model.predict(X_test)
```

```
mse = mean_squared_error(y_test, y_pred)
```

```
print(f'Mean Squared Error: {mse}')
```

```
print("Predicted relevance scores:", y_pred)
```

Output:-

```
>>> |  
===== RESTART: E:\Practical\IR\practical9a.py =====  
Mean Squared Error: 2.590136054421768  
Predicted relevance scores: [3.21428571 4.78571429 4.5      ]  
>>> |
```

- **Train the ranking model using labelled data and evaluate its effectiveness.**

Code:-

```
import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.ensemble import GradientBoostingRegressor
from sklearn.metrics import mean_squared_error, ndcg_score
data = {
    'doc_length': [100, 200, 150, 80, 120, 220, 90, 300, 250,
180],
    'keyword_freq': [10, 15, 8, 5, 9, 18, 4, 20, 14, 10],
    'avg_sentence_length': [10, 15, 12, 8, 9, 20, 7, 25, 18, 16],
    'relevance': [2, 3, 1, 0, 1, 3, 0, 3, 2, 2]
}
df = pd.DataFrame(data)
X = df[['doc_length', 'keyword_freq', 'avg_sentence_length']]
y = df['relevance']
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)
model = GradientBoostingRegressor(n_estimators=100,
learning_rate=0.1, max_depth=3)
```

```

model.fit(X_train, y_train)
y_pred = model.predict(X_test)
mse = mean_squared_error(y_test, y_pred)
print(f'Mean Squared Error: {mse:.4f}')
y_true = np.array([[0, 1, 2], [2, 1, 0]])
y_pred_reshaped = np.array([[0.1, 0.2, 0.7], [0.7, 0.2, 0.1]])
ndcg = ndcg_score(y_true, y_pred_reshaped)
print(f'NDCG Score: {ndcg:.4f}')
def mean_reciprocal_rank(y_true, y_pred):
    mrr = 0
    for true, pred in zip(y_true, y_pred):
        rank = np.argsort(pred)[::-1]
        for i, r in enumerate(rank):
            if true[r] > 0:
                mrr += 1 / (i + 1)
                break
    return mrr / len(y_true)
mrr_score = mean_reciprocal_rank(y_true, y_pred_reshaped)
print(f'MRR Score: {mrr_score:.4f}')

```

Output :-

```
===== RESTART: E:/Practical/IR/practical9b.py =====
```

```
Mean Squared Error: 0.4014
```

```
NDCG Score: 1.0000
```

```
MRR Score: 1.0000
```

Practical No.10

Aim:- Advanced Topics in Information Retrieval

- **Implement a text summarization algorithm (e.g., extractive or abstractive).**

Code:-

```
import nltk

from sklearn.feature_extraction.text import TfidfVectorizer
from nltk.tokenize import sent_tokenize
import heapq

nltk.download('punkt')

def summarize_text(text, summary_length=3):
    sentences = sent_tokenize(text)
    vectorizer = TfidfVectorizer(stop_words='english')
    tfidf_matrix = vectorizer.fit_transform(sentences)
    sentence_scores = {}
    for i, sentence in enumerate(sentences):
        sentence_scores[i] = tfidf_matrix[i].sum()
    ranked_sentences = heapq.nlargest(summary_length,
    sentence_scores, key=sentence_scores.get)
    ranked_sentences = sorted(ranked_sentences)
    summary = [sentences[i] for i in ranked_sentences]
```

```
return ' '.join(summary)
```

```
text = """
```

Natural language processing (NLP) is a subfield of linguistics, computer science, and artificial intelligence

concerned with the interactions between computers and human (natural) languages,

and, in particular, concerned with programming computers to fruitfully process large natural language data.

Challenges in natural language processing frequently involve natural language understanding, natural

language generation frequently from formal, machine-readable logical forms), connecting language and machine perception,

managing or analyzing large amounts of unstructured data (such as "big data"), and often statistical machine learning.

```
"""
```

```
summary = summarize_text(text)
```

```
print(summary)
```

```
nlTK.download('punkt_tab')
```

Output:-

```
>>>
===== RESTART: E:\Practical\IR\practical10a.py =====
[nltk_data] Downloading package punkt to
[nltk_data]   C:\Users\kc140\AppData\Roaming\nltk_data...
[nltk_data] Package punkt is already up-to-date!

Natural language processing (NLP) is a subfield of linguistics, computer science, and artificial intelligence
concerned with the interactions between computers and human (natural) languages,
and, in particular, concerned with programming computers to fruitfully process large natural language
data.

[nltk_data] Downloading package punkt_tab to
[nltk_data]   C:\Users\kc140\AppData\Roaming\nltk_data...
[nltk_data] Unzipping tokenizers\punkt_tab.zip.
>>>
```

- **Build a question-answering system using techniques such as information extraction**

Code:-

```
import nltk
import spacy
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.metrics.pairwise import cosine_similarity
from nltk.tokenize import sent_tokenize
import re
nltk.download('punkt')
def preprocess_text(text):
    return sent_tokenize(text)
def extract_relevant_sentences(text, question, top_n=3):
    sentences = preprocess_text(text)
    if not sentences:
        return []
    vectorizer = TfidfVectorizer(stop_words='english')
    tfidf_matrix = vectorizer.fit_transform([question] +
sentences)
    cosine_similarities = cosine_similarity(tfidf_matrix[0:1],
tfidf_matrix[1:]).flatten()
```



```

    relevant_sentences_idx = cosine_similarities.argsort()[-
top_n:][::-1]
    return [sentences[i] for i in relevant_sentences_idx]
def extract_entities(text):
    doc = nlp(text)
    entities = [(ent.text, ent.label_) for ent in doc.ents]
    return entities
def extract_answer_from_sentence(sentence, question):
    if "who" in question.lower():
        entities = extract_entities(sentence)
        for entity, label in entities:
            if label == "PERSON":
                return entity
    elif "what" in question.lower():
        words = re.findall(r'\b\w+\b', sentence)
        return " ".join(words[:5])
    return None
def answer_question(text, question, top_n=3):
    relevant_sentences = extract_relevant_sentences(text,
question, top_n)
    for sentence in relevant_sentences:

```

```
        answer = extract_answer_from_sentence(sentence,
question)
```

```
    if answer:
```

```
        return answer
```

```
    return "Sorry, I couldn't find an answer."
```

```
text = ""
```

Albert Einstein was born in Ulm, Germany, in 1879. He was a theoretical physicist who developed the theory of relativity.

Einstein is also known for his famous equation, $E = mc^2$. He won the Nobel Prize in Physics in 1921 for his discovery of the photoelectric effect. Einstein passed away in 1955 at the age of 76 in Princeton, New Jersey.

```
"""
```

```
questions = [
```

```
    "Who was Albert Einstein?",
```

```
    "What is  $E = mc^2$ ?",
```

```
    "When did Einstein pass away?"
```

```
]
```

```
nltk.download('punkt_tab')
```

```
nlp = spacy.load("en_core_web_sm")
```

```
for question in questions:
```

```
    print(f"Question: {question}")
```

```
print(f'Answer: {answer_question(text, question)}')  
print("-" * 40)
```

Output:-

```
>>> ===== RESTART: E:\Practical\IR\practical10b.py =====  
[nltk_data] Downloading package punkt to  
[nltk_data]   C:\Users\kc140\AppData\Roaming\nltk_data...  
[nltk_data] Package punkt is already up-to-date!  
[nltk_data] Downloading package punkt_tab to  
[nltk_data]   C:\Users\kc140\AppData\Roaming\nltk_data...  
[nltk_data] Package punkt_tab is already up-to-date!  
Question: Who was Albert Einstein?  
Answer: Albert Einstein  
  
-----  
Question: What is  $E = mc^2$ ?  
Answer: Einstein is also known for  
  
-----  
Question: When did Einstein pass away?  
Answer: Sorry, I couldn't find an answer.  
  
-----  
>>>
```